

Problem Statement

Build the linear regression model using scikit learn in boston data to predict 'Price' based on other dependent variable.

```
In [62]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import scipy.stats as stats
import sklearn
%matplotlib inline
```

```
In [63]: from sklearn.datasets import load_boston

boston = load_boston()
```

```
In [64]: boston.keys()
```

```
Out[64]: dict_keys(['data', 'target', 'feature_names', 'DESCR', 'filename'])
```

```
In [65]: boston.feature_names
```

```
Out[65]: array(['CRIM', 'ZN', 'INDUS', 'CHAS', 'NOX', 'RM', 'AGE', 'DIS', 'RAD',
               'TAX', 'PTRATIO', 'B', 'LSTAT'], dtype='<U7')
```

```
In [66]: print(boston.DESCR)
```

```
.. _boston_dataset:
```

```
Boston house prices dataset
-----
```

```
**Data Set Characteristics:**
```

```
:Number of Instances: 506
```

```
:Number of Attributes: 13 numeric/categorical predictive. Median Value (attribute 14) is usually the target.
```

```
:Attribute Information (in order):
```

- CRIM per capita crime rate by town
- ZN proportion of residential land zoned for lots over 25,000 sq.ft.
- INDUS proportion of non-retail business acres per town
- CHAS Charles River dummy variable (= 1 if tract bounds river; 0 otherwise)
- NOX nitric oxides concentration (parts per 10 million)
- RM average number of rooms per dwelling
- AGE proportion of owner-occupied units built prior to 1940
- DIS weighted distances to five Boston employment centres
- RAD index of accessibility to radial highways
- TAX full-value property-tax rate per \$10,000
- PTRATIO pupil-teacher ratio by town
- B $1000(Bk - 0.63)^2$ where Bk is the proportion of blacks by town
- LSTAT % lower status of the population
- MEDV Median value of owner-occupied homes in \$1000's

```
:Missing Attribute Values: None
```

```
:Creator: Harrison, D. and Rubinfeld, D.L.
```

```
In [67]: bos = pd.DataFrame(boston.data)
```

```
In [68]: bos.head()
```

```
Out[68]:
```

	0	1	2	3	4	5	6	7	8	9	10	11	12
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	15.3	396.90	4.98
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	17.8	396.90	9.14
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	17.8	392.83	4.03
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	18.7	394.63	2.94
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	18.7	396.90	5.33

```
In [69]: bos.shape
```

```
Out[69]: (506, 13)
```

```
In [70]: bos.columns = boston.feature_names
```

```
In [71]: bos.head()
```

```
Out[71]:
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LSTAT
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	15.3	396.90	4.98
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	17.8	396.90	9.14
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	17.8	392.83	4.03
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	18.7	394.63	2.94
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	18.7	396.90	5.33

```
In [72]: bos['Price'] = boston.target
```

```
In [73]: bos.head()
```

```
Out[73]:
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LSTAT	Price
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	15.3	396.90	4.98	24.0
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	17.8	396.90	9.14	21.6
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	17.8	392.83	4.03	34.7
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	18.7	394.63	2.94	33.4
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	18.7	396.90	5.33	36.2

```
In [76]: bos.shape
```

```
Out[76]: (506, 14)
```

```
In [74]: X = bos.iloc[:, :-1].values
```

```
In [77]: X.shape
```

```
Out[77]: (506, 13)
```

```
In [78]: y = bos.iloc[:, 13].values
```

```
In [79]: y.shape
```

```
Out[79]: (506,)
```

```
In [105]: from sklearn.model_selection import train_test_split  
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2, random_state = 0)
```

```
In [106]: print('X_train: ', X_train.shape)  
print('X_test: ', X_test.shape)  
print('Y_train: ', y_train.shape)  
print('Y_test: ', y_test.shape)
```

```
X_train: (404, 13)  
X_test: (102, 13)  
Y_train: (404,)  
Y_test: (102,)
```

```
In [107]: from sklearn.linear_model import LinearRegression  
  
lm = LinearRegression()  
lm.fit(X_train, y_train)
```

```
Out[107]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None,  
normalize=False)
```

```
In [108]: y_pred = lm.predict(X_test)
```

```
In [134]: ActualvsPred = pd.DataFrame({'Actual': y_test, 'Predicted': y_pred})
```

```
In [135]: print(ActualvsPred.head())
```

	Actual	Predicted
0	22.6	24.889638
1	50.0	23.721411
2	23.0	29.364999
3	8.3	12.122386
4	21.2	21.443823

```
In [139]: print(ActualvsPred.tail())
```

	Actual	Predicted
97	24.7	25.442171
98	14.1	15.571783
99	18.7	17.937195
100	28.1	25.305888
101	19.8	22.373233

```
In [138]: plt.scatter(y_test,y_pred, color = 'blue')
plt.xlabel('Actual Prices')
plt.ylabel('Predicted Prices')
plt.title('Actual Prices vs Prices Predicted')
plt.show
```

```
Out[138]: <function matplotlib.pyplot.show(*args, **kw)>
```



```
In [125]: print(lm.coef_)

[-1.19443447e-01  4.47799511e-02  5.48526168e-03  2.34080361e+00
 -1.61236043e+01  3.70870901e+00 -3.12108178e-03 -1.38639737e+00
  2.44178327e-01 -1.09896366e-02 -1.04592119e+00  8.11010693e-03
 -4.92792725e-01]
```

```
In [140]: print(lm.intercept_)

38.091694926302004
```

```
In [140]: print(lm.intercept_)

38.091694926302004
```

```
In [141]: from sklearn.metrics import r2_score

score = r2_score(y_test,y_pred)
print('R Sqaured Score of the Test data is: ', score)

R Sqaured Score of the Test data is:  0.5892223849182534
```

```
In [143]: print('Mean Absolute Error (MAE) of Test data is: ',metrics.mean_absolute_error(y_test,y_pred))
print('Mean Squared Error (MSE) of Test data is: ',metrics.mean_squared_error(y_test,y_pred))
print('Root Mean Squared Error (RMSE) of Test data is: ',np.sqrt(metrics.mean_squared_error(y_test,y_pred)))

Mean Absolute Error (MAE) of Test data is:  3.8429092204444912
Mean Squared Error (MSE) of Test data is:  33.44897999767632
Root Mean Squared Error (RMSE) of Test data is:  5.783509315085117
```

```
In [144]: ## The R2 Score of 0.58922 is Low, therefore our model is not very accurate at predicting the
## house prices based on the data provided.
```