

# Project 2

## Predicting Players Rating

The dataset comes in the form of an SQL database and contains statistics of about 25,000 football matches, from the top football league of 11 European Countries. It covers seasons from 2008 to 2016 and contains match statistics (i.e: scores, corners, fouls etc...) as well as the team formations, with player names and a pair of coordinates to indicate their position on the pitch.

In this project you are going to predict the overall rating of soccer player based on their attributes such as 'crossing', 'finishing etc.

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns
import sqlite3
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split
from sklearn import metrics
from sklearn.metrics import r2_score
from math import sqrt
import statsmodels.formula.api as smf
```

```
In [2]: # Create the connection to the dB
# Read the Player Attributes table data into a dataframe

conn = sqlite3.connect('database.sqlite')
df = pd.read_sql_query("SELECT * FROM Player_Attributes", conn)
```

In [3]: *# View the first view rows of the data*

```
df.head()
```

Out[3]:

	id	player_fifa_api_id	player_api_id	date	overall_rating	potential	preferred_foot	attacking_work_rate	defensive_work_rate	crossin
0	1	218353	505942	2016-02-18 00:00:00	67.0	71.0	right	medium	medium	49
1	2	218353	505942	2015-11-19 00:00:00	67.0	71.0	right	medium	medium	49
2	3	218353	505942	2015-09-21 00:00:00	62.0	66.0	right	medium	medium	49
3	4	218353	505942	2015-03-20 00:00:00	61.0	65.0	right	medium	medium	48
4	5	218353	505942	2007-02-22 00:00:00	61.0	65.0	right	medium	medium	48

5 rows × 42 columns



In [4]: *# The number of rows and columns in our original dataset*

```
df.shape
```

Out[4]: (183978, 42)

In [5]: *# Information about our columns and their data type*

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 183978 entries, 0 to 183977
Data columns (total 42 columns):
id                183978 non-null int64
player_fifa_api_id  183978 non-null int64
player_api_id     183978 non-null int64
date              183978 non-null object
overall_rating     183142 non-null float64
potential         183142 non-null float64
preferred_foot     183142 non-null object
attacking_work_rate 180748 non-null object
defensive_work_rate 183142 non-null object
crossing           183142 non-null float64
finishing          183142 non-null float64
heading_accuracy   183142 non-null float64
short_passing      183142 non-null float64
volleys            181265 non-null float64
dribbling          183142 non-null float64
curve              181265 non-null float64
free_kick_accuracy 183142 non-null float64
long_passing       183142 non-null float64
ball_control       183142 non-null float64
acceleration       183142 non-null float64
sprint_speed       183142 non-null float64
agility            181265 non-null float64
reactions          183142 non-null float64
balance            181265 non-null float64
shot_power         183142 non-null float64
jumping            181265 non-null float64
stamina            183142 non-null float64
strength           183142 non-null float64
long_shots         183142 non-null float64
aggression         183142 non-null float64
interceptions      183142 non-null float64
positioning        183142 non-null float64
vision             181265 non-null float64
penalties          183142 non-null float64
marking            183142 non-null float64
standing_tackle    183142 non-null float64
```

```
sliding_tackle      181265 non-null float64
gk_diving            183142 non-null float64
gk_handling          183142 non-null float64
gk_kicking           183142 non-null float64
gk_positioning       183142 non-null float64
gk_reflexes          183142 non-null float64
dtypes: float64(35), int64(3), object(4)
memory usage: 59.0+ MB
```

## Exploratory Data Analysis

### Data Cleaning

#### Model Evaluation without Categorical Variables

```
In [6]: # Drop the columns id, player_fifa_api_id, player_api_id, date since these features do not
        # significantly contribute to our model.

        # Also remove columns with categorical values preferred_foot, attacking_work_rate, and defensive_work_rate
        # to see how our model compares.

        df.drop(['id', 'player_fifa_api_id', 'player_api_id', 'date', 'preferred_foot', 'attacking_work_rate',
                  'defensive_work_rate'], axis=1, inplace=True)
```

```
In [7]: df.shape
```

```
Out[7]: (183978, 35)
```

In [8]: `df.head(10)`

Out[8]:

	overall_rating	potential	crossing	finishing	heading_accuracy	short_passing	volleys	dribbling	curve	free_kick_accuracy	...	vision
0	67.0	71.0	49.0	44.0	71.0	61.0	44.0	51.0	45.0	39.0	...	54.0
1	67.0	71.0	49.0	44.0	71.0	61.0	44.0	51.0	45.0	39.0	...	54.0
2	62.0	66.0	49.0	44.0	71.0	61.0	44.0	51.0	45.0	39.0	...	54.0
3	61.0	65.0	48.0	43.0	70.0	60.0	43.0	50.0	44.0	38.0	...	53.0
4	61.0	65.0	48.0	43.0	70.0	60.0	43.0	50.0	44.0	38.0	...	53.0
5	74.0	76.0	80.0	53.0	58.0	71.0	40.0	73.0	70.0	69.0	...	66.0
6	74.0	76.0	80.0	53.0	58.0	71.0	32.0	73.0	70.0	69.0	...	66.0
7	73.0	75.0	79.0	52.0	57.0	70.0	29.0	71.0	68.0	69.0	...	65.0
8	73.0	75.0	79.0	51.0	57.0	70.0	29.0	71.0	68.0	69.0	...	65.0
9	73.0	75.0	79.0	51.0	57.0	70.0	29.0	71.0	68.0	69.0	...	65.0

10 rows × 35 columns



```
In [9]: # Count the number of number of null values in the columns  
df.isnull().sum(axis=0)
```

```
Out[9]: overall_rating      836  
        potential          836  
        crossing           836  
        finishing          836  
        heading_accuracy   836  
        short_passing      836  
        volleys            2713  
        dribbling          836  
        curve              2713  
        free_kick_accuracy  836  
        long_passing       836  
        ball_control       836  
        acceleration       836  
        sprint_speed       836  
        agility            2713  
        reactions          836  
        balance            2713  
        shot_power         836  
        jumping            2713  
        stamina            836  
        strength           836  
        long_shots         836  
        aggression         836  
        interceptions      836  
        positioning       836  
        vision             2713  
        penalties         836  
        marking           836  
        standing_tackle    836  
        sliding_tackle     2713  
        gk_diving          836  
        gk_handling        836  
        gk_kicking         836  
        gk_positioning     836  
        gk_reflexes        836  
        dtype: int64
```

```
In [10]: # Since my data only contains numerical values I will fill the null values with mean  
df.fillna(df.mean(), inplace=True)
```

```
In [11]: df.isnull().sum(axis=0)
```

```
Out[11]: overall_rating      0
         potential          0
         crossing            0
         finishing           0
         heading_accuracy    0
         short_passing       0
         volleys             0
         dribbling           0
         curve               0
         free_kick_accuracy  0
         long_passing        0
         ball_control        0
         acceleration        0
         sprint_speed        0
         agility             0
         reactions           0
         balance             0
         shot_power          0
         jumping             0
         stamina             0
         strength            0
         long_shots          0
         aggression          0
         interceptions        0
         positioning         0
         vision              0
         penalties           0
         marking             0
         standing_tackle     0
         sliding_tackle      0
         gk_diving           0
         gk_handling         0
         gk_kicking          0
         gk_positioning      0
         gk_reflexes         0
         dtype: int64
```

**Since there is close to 184 thousand rows of data take a random sample to do some EDA to see if there is correlation between the features and the overall rating of the player**



```
In [12]: # Take a random sample of a thousand rows of data to do Exploratory Data Analysis
df_subset = df.sample(n=1000, random_state=0)
```

```
In [13]: df_subset.shape
```


```
Out[13]: (1000, 35)
```

```
In [14]: df_subset.describe()
```

```
Out[14]:
```

	overall_rating	potential	crossing	finishing	heading_accuracy	short_passing	volleys	dribbling	curve	fr
<b>count</b>	1000.000000	1000.000000	1000.000000	1000.000000	1000.000000	1000.000000	1000.000000	1000.000000	1000.000000	
<b>mean</b>	68.313600	73.272762	55.568521	49.630526	57.633596	62.934578	49.624153	59.326051	53.110622	
<b>std</b>	7.040126	6.555597	16.629112	18.481994	16.103922	13.443303	17.728343	16.899600	17.453120	
<b>min</b>	39.000000	51.000000	6.000000	7.000000	7.000000	12.000000	3.000000	6.000000	6.000000	
<b>25%</b>	64.000000	69.000000	47.000000	35.000000	50.000000	58.000000	37.000000	53.000000	42.000000	
<b>50%</b>	69.000000	74.000000	59.000000	51.000000	60.000000	65.000000	52.000000	63.000000	55.500000	
<b>75%</b>	73.000000	78.000000	68.000000	64.000000	69.000000	72.000000	64.000000	71.000000	66.000000	
<b>max</b>	89.000000	92.000000	86.000000	95.000000	94.000000	93.000000	91.000000	92.000000	89.000000	


8 rows × 35 columns



In [15]: *# See the distribution of overall rating data.*

```
plt.figure(figsize=(15,8))  
df_subset['overall_rating'].hist(bins=20, color='red', alpha=0.5)  
plt.show
```


Out[15]: <function matplotlib.pyplot.show(\*args, \*\*kw)>



In [16]: *# See the distribution of agility data.*

```
plt.figure(figsize=(15,8))  
df_subset['agility'].hist(bins=20, color='blue', alpha=0.5)  
plt.show
```


Out[16]: <function matplotlib.pyplot.show(\*args, \*\*kw)>



In [17]: *# See if there is a relationship between agility and the overall rating of the player.*


```
plt.figure(figsize=(15,8))
sns.set_style('whitegrid')
sns.scatterplot(x='agility', y='overall_rating', data=df_subset, hue='overall_rating', palette='Reds')
plt.show
```

Out[17]: <function matplotlib.pyplot.show(\*args, \*\*kw)>




```
In [18]: # See if there is a relationship between dribbling and the overall rating of the player.  
  
plt.figure(figsize=(15,8))  
sns.set_style('whitegrid')  
sns.scatterplot(x='dribbling', y='overall_rating', data=df_subset, hue='overall_rating', palette='rainbow')  
plt.show
```

```
Out[18]: <function matplotlib.pyplot.show(*args, **kw)>
```



```
In [19]: # See if there is a relationship between free kick accuracy and the overall rating of the player.  
  
plt.figure(figsize=(15,8))  
sns.set_style('whitegrid')  
sns.scatterplot(x='free_kick_accuracy', y='overall_rating', data=df_subset, hue='overall_rating', palette='Greens',  
plt.show
```

```
Out[19]: <function matplotlib.pyplot.show(*args, **kw)>
```



In [20]: *# See if there is a relationship between a players vision and the overall rating.*


```
plt.figure(figsize=(15,8))
sns.set_style('whitegrid')
sns.scatterplot(x='vision', y='overall_rating', data=df_subset, hue='overall_rating', palette='Blues')
plt.show
```

Out[20]: <function matplotlib.pyplot.show(\*args, \*\*kw)>



```
In [21]: # See if there is a relationship between a players marking ability and the overall rating.  
  
plt.figure(figsize=(15,8))  
sns.set_style('whitegrid')  
sns.scatterplot(x='potential', y='overall_rating', data=df_subset, hue='overall_rating', palette='Oranges')  
plt.show
```

```
Out[21]: <function matplotlib.pyplot.show(*args, **kw)>
```



## Building Linear Regression Model



## Split the data in to Training and Test set

```
In [22]: X_train, X_test, y_train, y_test = train_test_split(df.drop('overall_rating', axis=1), df['overall_rating'],
test_size = 0.30, random_state = 101)
```

```
In [23]: X_train.head()
```

Out[23]:

	potential	crossing	finishing	heading_accuracy	short_passing	volleys	dribbling	curve	free_kick_accuracy	long_passing	...	\
<b>93681</b>	75.0	62.0	44.0	53.0	62.0	55.0	65.0	58.0	68.0	54.0	...	
<b>116493</b>	83.0	39.0	35.0	86.0	64.0	43.0	46.0	33.0	27.0	62.0	...	
<b>17520</b>	73.0	65.0	71.0	71.0	64.0	65.0	69.0	59.0	56.0	54.0	...	
<b>97796</b>	68.0	53.0	68.0	71.0	63.0	67.0	61.0	52.0	65.0	55.0	...	
<b>124921</b>	67.0	57.0	49.0	63.0	67.0	59.0	66.0	58.0	56.0	64.0	...	

5 rows × 34 columns

```
In [24]: y_train.head()
```

```
Out[24]: 93681    64.0
116493    81.0
17520     71.0
97796     68.0
124921    67.0
Name: overall_rating, dtype: float64
```

```
In [25]: # Print the shape of train and test data
```

```
print('X_train: ', X_train.shape)
print('X_test: ', X_test.shape)
print('y_train: ', y_train.shape)
print('y_test: ', y_test.shape)
```

```
X_train: (128784, 34)
X_test: (55194, 34)
y_train: (128784,)
y_test: (55194,)
```

## Create and Train the Linear Regression Model

```
In [26]: lm = LinearRegression()  
lm.fit(X_train, y_train)
```

```
Out[26]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None,  
normalize=False)
```

## Evaluate the model by checking out it's coefficients

```
In [27]: # Find the intercept  
print(lm.intercept_)
```

```
-3.7138600642612545
```

```
In [28]: # Find the Coefficient of X train data
coeff_df = pd.DataFrame(lm.coef_,X_train.columns,columns=['Coefficient'])
print(coeff_df)
```

	Coefficient
potential	0.381386
crossing	0.021691
finishing	0.012632
heading_accuracy	0.069519
short_passing	0.049780
volleys	0.002269
dribbling	-0.012380
curve	0.010761
free_kick_accuracy	0.014515
long_passing	0.006448
ball_control	0.135916
acceleration	0.005902
sprint_speed	0.009651
agility	-0.008306
reactions	0.208145
balance	0.007479
shot_power	0.015337
jumping	0.016442
stamina	-0.005122
strength	0.061832
long_shots	-0.011895
aggression	0.020297
interceptions	0.013100
positioning	-0.010003
vision	-0.002524
penalties	0.013118
marking	0.034321
standing_tackle	0.002232
sliding_tackle	-0.029936
gk_diving	0.166042
gk_handling	0.031718
gk_kicking	-0.032619
gk_positioning	0.054697
gk_reflexes	0.022835

## Make the Predictions

```
In [29]: y_pred = lm.predict(X_test)
```

```
In [30]: # Compare the Actual Overall Rating to the Predicted Overall Rating  
  
ActualvsPred = pd.DataFrame({'Actual': y_test, 'Predicted': y_pred})
```

```
In [31]: print(ActualvsPred.head(20))
```

	Actual	Predicted
89795	81.0	76.500393
145987	72.0	68.447011
81345	65.0	64.701619
40399	76.0	73.147422
143301	76.0	72.540214
134213	70.0	70.926820
183666	74.0	76.433557
89101	57.0	56.650473
48414	71.0	70.949247
124003	63.0	62.383194
179632	70.0	70.053771
101291	64.0	63.346655
95221	78.0	79.922389
30845	69.0	61.572655
4734	74.0	69.891753
53097	74.0	73.911948
128101	68.0	68.187802
12918	69.0	64.483406
50195	59.0	67.410854
162266	63.0	62.247938

```
In [32]: # Calculate the R Squared value of the Actual Overall Rating to the Predicted Overall Rating  
score = r2_score(y_test,y_pred)  
print('R Sqaured Score of the Test data is: ', score)
```

R Sqaured Score of the Test data is: 0.8424268878510373

```
In [33]: print('Mean Absolute Error (MAE) of Test data is: ',metrics.mean_absolute_error(y_test,y_pred))
print('Mean Squared Error (MSE) of Test data is: ',metrics.mean_squared_error(y_test,y_pred))
print('Root Mean Squared Error (RMSE) of Test data is: ',np.sqrt(metrics.mean_squared_error(y_test,y_pred)))
```

Mean Absolute Error (MAE) of Test data is: 2.124357660885702  
Mean Squared Error (MSE) of Test data is: 7.7783783388997545  
Root Mean Squared Error (RMSE) of Test data is: 2.788974424210404

```
In [34]: # The R2 Score of 0.8424, therefore our model can predict the overall rating of the players with
# approximately 84% accuracy based on the player attributes.
```

## Model Evaluation with Categorical Variables

```
In [35]: conn = sqlite3.connect('database.sqlite')
df1 = pd.read_sql_query("SELECT * FROM Player_Attributes", conn)
```

```
In [36]: df1.drop(['id', 'player_fifa_api_id', 'player_api_id', 'date'], axis=1, inplace=True)
```

```
In [37]: df1.head()
```

Out[37]:

	overall_rating	potential	preferred_foot	attacking_work_rate	defensive_work_rate	crossing	finishing	heading_accuracy	short_passing
0	67.0	71.0	right	medium	medium	49.0	44.0	71.0	61.0
1	67.0	71.0	right	medium	medium	49.0	44.0	71.0	61.0
2	62.0	66.0	right	medium	medium	49.0	44.0	71.0	61.0
3	61.0	65.0	right	medium	medium	48.0	43.0	70.0	60.0
4	61.0	65.0	right	medium	medium	48.0	43.0	70.0	60.0

5 rows × 38 columns



```
In [38]: df1.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 183978 entries, 0 to 183977
Data columns (total 38 columns):
overall_rating      183142 non-null float64
potential           183142 non-null float64
preferred_foot      183142 non-null object
attacking_work_rate 180748 non-null object
defensive_work_rate 183142 non-null object
crossing            183142 non-null float64
finishing           183142 non-null float64
heading_accuracy    183142 non-null float64
short_passing       183142 non-null float64
volleys            181265 non-null float64
dribbling           183142 non-null float64
curve              181265 non-null float64
free_kick_accuracy  183142 non-null float64
long_passing        183142 non-null float64
ball_control        183142 non-null float64
acceleration        183142 non-null float64
sprint_speed        183142 non-null float64
agility             181265 non-null float64
reactions           183142 non-null float64
balance             181265 non-null float64
shot_power          183142 non-null float64
jumping             181265 non-null float64
stamina             183142 non-null float64
strength            183142 non-null float64
long_shots          183142 non-null float64
aggression          183142 non-null float64
interceptions       183142 non-null float64
positioning         183142 non-null float64
vision              181265 non-null float64
penalties           183142 non-null float64
marking             183142 non-null float64
standing_tackle     183142 non-null float64
sliding_tackle      181265 non-null float64
gk_diving           183142 non-null float64
gk_handling         183142 non-null float64
gk_kicking          183142 non-null float64
gk_positioning      183142 non-null float64
gk_reflexes         183142 non-null float64
```

```
dtypes: float64(35), object(3)  
memory usage: 53.3+ MB
```

```
In [39]: df1.isnull().sum(axis=0)
```

```
Out[39]: overall_rating      836  
         potential          836  
         preferred_foot      836  
         attacking_work_rate 3230  
         defensive_work_rate  836  
         crossing            836  
         finishing           836  
         heading_accuracy    836  
         short_passing       836  
         volleys             2713  
         dribbling           836  
         curve               2713  
         free_kick_accuracy   836  
         long_passing        836  
         ball_control         836  
         acceleration        836  
         sprint_speed        836  
         agility             2713  
         reactions           836  
         balance             2713  
         shot_power          836  
         jumping             2713  
         stamina             836  
         strength            836  
         long_shots          836  
         aggression          836  
         interceptions        836  
         positioning         836  
         vision              2713  
         penalties           836  
         marking             836  
         standing_tackle     836  
         sliding_tackle      2713  
         gk_diving           836  
         gk_handling          836  
         gk_kicking           836  
         gk_positioning      836  
         gk_reflexes          836  
         dtype: int64
```



```
In [40]: pd.value_counts(df1['preferred_foot'])
```

```
Out[40]: right    138409  
left      44733  
Name: preferred_foot, dtype: int64
```

```
In [43]: # See what other values are in the attacking work rate column aside from the null values
```

```
pd.value_counts(df1['attacking_work_rate'])
```

```
Out[43]: medium    125070  
high      42823  
low       8569  
None      3639  
norm       348  
y          106  
le         104  
stoc        89  
Name: attacking_work_rate, dtype: int64
```

```
In [44]: # See what other values are in the defensive work rate column aside from the null values  
pd.value_counts(df1['defensive_work_rate'])
```

```
Out[44]: medium      130846  
         high        27041  
         low         18432  
         _0          2394  
         o           1550  
         1            441  
         ormal        348  
         2            342  
         3            258  
         5            234  
         7            217  
         0            197  
         6            197  
         9            152  
         4            116  
         es           106  
         ean           104  
         tocky         89  
         8             78  
Name: defensive_work_rate, dtype: int64
```

```
In [45]: # Create a function to replace null values in the preferred foot column with the preferred  
         # foot as right since most players are right footed.
```

```
def impute_preferred_foot (cols):  
    preferred_foot = cols[0]  
  
    if pd.isnull(preferred_foot):  
        return 'right'  
    else:  
        return preferred_foot
```

```
In [46]: # Apply the function to replace the null values in the preferred foot column
```

```
df1['preferred_foot'] = df1[['preferred_foot']].apply(impute_preferred_foot, axis=1)
```

```
In [47]: df1.isnull().sum(axis=0)
```

```
Out[47]: overall_rating      836  
         potential          836  
         preferred_foot      0  
         attacking_work_rate 3230  
         defensive_work_rate 836  
         crossing            836  
         finishing           836  
         heading_accuracy    836  
         short_passing       836  
         volleys             2713  
         dribbling           836  
         curve               2713  
         free_kick_accuracy   836  
         long_passing        836  
         ball_control         836  
         acceleration        836  
         sprint_speed        836  
         agility             2713  
         reactions           836  
         balance             2713  
         shot_power          836  
         jumping             2713  
         stamina             836  
         strength            836  
         long_shots          836  
         aggression          836  
         interceptions       836  
         positioning         836  
         vision              2713  
         penalties           836  
         marking             836  
         standing_tackle     836  
         sliding_tackle      2713  
         gk_diving           836  
         gk_handling         836  
         gk_kicking          836  
         gk_positioning      836  
         gk_reflexes         836  
         dtype: int64
```

```
In [48]: pd.value_counts(df1['preferred_foot'])
```

```
Out[48]: right    139245  
left      44733  
Name: preferred_foot, dtype: int64
```

```
In [49]: # Create a function to replace None, norm and null values in the attacking work rate column with  
# medium since most players are classified with medium attacking rate. Replace y, le, and stoc.  
  
def impute_attacking_rate (cols):  
    attacking_work_rate = cols[0]  
  
    if pd.isnull(attacking_work_rate):  
        return 'medium'  
    elif attacking_work_rate == 'None' or attacking_work_rate == 'norm':  
        return 'medium'  
    elif attacking_work_rate == 'y' or attacking_work_rate == 'le' or attacking_work_rate == 'stoc':  
        return 'low'  
  
    else:  
        return attacking_work_rate
```

```
In [50]: # Apply the function to replace the None, norm, and null values in the attacking work rate column  
  
df1['attacking_work_rate'] = df1[['attacking_work_rate']].apply(impute_attacking_rate, axis=1)
```

```
In [51]: pd.value_counts(df1['attacking_work_rate'])
```

```
Out[51]: medium    132287  
high      42823  
low       8868  
Name: attacking_work_rate, dtype: int64
```

```
In [52]: # Create a function to replace null and other values that is not high, medium or low in the defensive work rate  
# either high, medium or low.
```

```
def impute_defensive_rate (cols):  
    defensive_work_rate = cols[0]  
  
    if pd.isnull(defensive_work_rate):  
        return 'medium'  
    elif defensive_work_rate in ('_0', 'o', '0', '1', '2', '3'):  
        return 'low'  
    elif defensive_work_rate in ('ormal', 'es', 'ean', 'tocky', '4', '5', '6'):  
        return 'medium'  
    elif defensive_work_rate in ('7', '8', '9'):  
        return 'high'  
  
    else:  
        return defensive_work_rate
```

```
In [53]: # Apply the function to replace the values in the defensive work rate column
```

```
df1['defensive_work_rate'] = df1[['defensive_work_rate']].apply(impute_defensive_rate, axis=1)
```

```
In [54]: pd.value_counts(df1['defensive_work_rate'])
```

```
Out[54]: medium    132876  
         high      27488  
         low       23614  
         Name: defensive_work_rate, dtype: int64
```

```
In [57]: # Replace all the null values in the columns with numerical values with mean values
```

```
df1.fillna(df.mean(), inplace=True)
```

```
In [58]: df1.isnull().sum(axis=0)
```

```
Out[58]: overall_rating      0
         potential          0
         preferred_foot      0
         attacking_work_rate  0
         defensive_work_rate  0
         crossing            0
         finishing           0
         heading_accuracy    0
         short_passing       0
         volleys             0
         dribbling           0
         curve               0
         free_kick_accuracy  0
         long_passing        0
         ball_control        0
         acceleration        0
         sprint_speed        0
         agility             0
         reactions           0
         balance             0
         shot_power          0
         jumping             0
         stamina             0
         strength            0
         long_shots          0
         aggression          0
         interceptions       0
         positioning         0
         vision              0
         penalties           0
         marking             0
         standing_tackle     0
         sliding_tackle      0
         gk_diving           0
         gk_handling         0
         gk_kicking          0
         gk_positioning      0
         gk_reflexes         0
         dtype: int64
```

```
In [59]: # We will convert the preferred foot, attacking work rate and defensive work rate columns  
# to dummy variable (LabelEncoding) and drop one column (OneHotEncoding)  
# This will create a new dataframe for each feature
```

```
pref_foot = pd.get_dummies(df1['preferred_foot'], drop_first=True)  
attack_rate = pd.get_dummies(df1['attacking_work_rate'], drop_first=True)  
def_rate = pd.get_dummies(df1['defensive_work_rate'], drop_first=True)
```

```
In [60]: pref_foot.head()
```

Out[60]:

	right
0	1
1	1
2	1
3	1
4	1

```
In [61]: attack_rate.head()
```

Out[61]:

	low	medium
0	0	1
1	0	1
2	0	1
3	0	1
4	0	1

```
In [62]: def_rate.head()
```

```
Out[62]:
```

	low	medium
0	0	1
1	0	1
2	0	1
3	0	1
4	0	1

```
In [64]: # We will drop the preferred foot, attacking work rate, and defensive work rate columns from the  
# original dataset since we created the dummy variable.  
df1.drop(['preferred_foot', 'attacking_work_rate', 'defensive_work_rate'], axis=1, inplace=True)
```

```
In [65]: # We will concatenate the pref_foot, attack_rate, and def_rate dummy variables to our dataset.  
df1 = pd.concat([df1,pref_foot, attack_rate, def_rate],axis=1)
```

```
In [66]: df1.head()
```

```
Out[66]:
```

t_passing	volleys	dribbling	curve	free_kick_accuracy	...	gk_diving	gk_handling	gk_kicking	gk_positioning	gk_reflexes	right	low	me
61.0	44.0	51.0	45.0	39.0	...	6.0	11.0	10.0	8.0	8.0	1	0	
61.0	44.0	51.0	45.0	39.0	...	6.0	11.0	10.0	8.0	8.0	1	0	
61.0	44.0	51.0	45.0	39.0	...	6.0	11.0	10.0	8.0	8.0	1	0	
60.0	43.0	50.0	44.0	38.0	...	5.0	10.0	9.0	7.0	7.0	1	0	
60.0	43.0	50.0	44.0	38.0	...	5.0	10.0	9.0	7.0	7.0	1	0	

## Building Linear Regression Model



## Split the data in to Training and Test set

```
In [74]: X_train1, X_test1, y_train1, y_test1 = train_test_split(df1.drop('overall_rating', axis=1), df1['overall_rating'],
                                                                test_size = 0.30, random_state = 101)
```

## Create and Train the Linear Regression Model

```
In [75]: lm1 = LinearRegression()
lm1.fit(X_train1, y_train1)
```

```
Out[75]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None,
                           normalize=False)
```

## Evaluate the model by checking out it's coefficients

```
In [76]: # Find the intercept
print(lm1.intercept_)
```

```
-3.385317880706424
```

```
In [77]: # Find the Coefficient of X train data
coeff_df1 = pd.DataFrame(lm1.coef_,X_train.columns,columns=['Coefficient'])
print(coeff_df1)
```

	Coefficient
potential	0.379404
crossing	0.022461
finishing	0.013157
heading_accuracy	0.068397
short_passing	0.051094
volleys	0.001255
dribbling	-0.010241
curve	0.011260
free_kick_accuracy	0.014685
long_passing	0.006015
ball_control	0.134659
acceleration	0.005330
sprint_speed	0.011517
agility	-0.008333
reactions	0.206745
balance	0.006603
shot_power	0.015242
jumping	0.016285
stamina	-0.003735
strength	0.060123
long_shots	-0.012203
aggression	0.019422
interceptions	0.012227
positioning	-0.008913
vision	-0.002528
penalties	0.012153
marking	0.033513
standing_tackle	0.002884
sliding_tackle	-0.028418
gk_diving	0.168623
gk_handling	0.032292
gk_kicking	-0.034173
gk_positioning	0.056890
gk_reflexes	0.023825
right	-0.012523
low	1.201846
medium	-0.108176

low	0.237049
medium	-0.280066

## Make the Predictions

```
In [78]: y_pred1 = lm1.predict(X_test1)
```

```
In [79]: # Compare the Actual Overall Rating to the Predicted Overall Rating  
  
ActualvsPred1 = pd.DataFrame({'Actual': y_test1, 'Predicted': y_pred1})
```

```
In [80]: print(ActualvsPred1.head(20))
```

	Actual	Predicted
89795	81.0	76.111882
145987	72.0	68.868899
81345	65.0	64.514367
40399	76.0	74.088600
143301	76.0	72.186087
134213	70.0	71.078865
183666	74.0	77.415047
89101	57.0	56.946604
48414	71.0	70.983164
124003	63.0	62.456015
179632	70.0	70.228510
101291	64.0	63.799896
95221	78.0	79.896610
30845	69.0	61.548103
4734	74.0	69.939717
53097	74.0	73.631112
128101	68.0	68.003258
12918	69.0	65.502644
50195	59.0	68.702345
162266	63.0	62.179165

In [81]: *# Calculate the R Squared value of the Actual Overall Rating to the Predicted Overall Rating*

```
score1 = r2_score(y_test1,y_pred1)
print('R Sqaured Score of the Test data is: ', score1)
```

R Sqaured Score of the Test data is: 0.8445423047381513

In [82]: 

```
print('Mean Absolute Error (MAE) of Test data is: ',metrics.mean_absolute_error(y_test1,y_pred1))
print('Mean Squared Error (MSE) of Test data is: ',metrics.mean_squared_error(y_test1,y_pred1))
print('Root Mean Squared Error (RMSE) of Test data is: ',np.sqrt(metrics.mean_squared_error(y_test1,y_pred1)))
```

Mean Absolute Error (MAE) of Test data is: 2.111661475923508

Mean Squared Error (MSE) of Test data is: 7.673953715510228

Root Mean Squared Error (RMSE) of Test data is: 2.770190194826021

In [83]: *# The R2 Score of 0.8445, therefore our model can predict the overall rating of the players with  
# approximately 84% accuracy based on the player attributes. The R2 score did not change significantly therefore  
# the categorical columns did not impact my model.*

In [ ]: