

```
In [125]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import r2_score, mean_absolute_error, mean_squared_error
from sklearn.preprocessing import StandardScaler, LabelEncoder
```

```
In [84]: from sklearn.datasets import load_boston

boston = load_boston()
```

```
In [85]: boston.feature_names
```

```
Out[85]: array(['CRIM', 'ZN', 'INDUS', 'CHAS', 'NOX', 'RM', 'AGE', 'DIS', 'RAD',
               'TAX', 'PTRATIO', 'B', 'LSTAT'], dtype='<U7')
```

```
In [86]: print(boston.DESCR)
```

```
.. _boston_dataset:
```

```
Boston house prices dataset
```

```
-----
```

```
**Data Set Characteristics:**
```

```
:Number of Instances: 506
```

```
:Number of Attributes: 13 numeric/categorical predictive. Median Value (attribute 14) is usually the target.
```

```
:Attribute Information (in order):
```

- CRIM per capita crime rate by town
- ZN proportion of residential land zoned for lots over 25,000 sq.ft.
- INDUS proportion of non-retail business acres per town
- CHAS Charles River dummy variable (= 1 if tract bounds river; 0 otherwise)
- NOX nitric oxides concentration (parts per 10 million)
- RM average number of rooms per dwelling
- AGE proportion of owner-occupied units built prior to 1940
- DIS weighted distances to five Boston employment centres
- RAD index of accessibility to radial highways
- TAX full-value property-tax rate per \$10,000
- PTRATIO pupil-teacher ratio by town
- B $1000(B_k - 0.63)^2$ where B_k is the proportion of blacks by town
- LSTAT % lower status of the population
- MEDV Median value of owner-occupied homes in \$1000's

```
:Missing Attribute Values: None
```

```
:Creator: Harrison, D. and Rubinfeld, D.L.
```

This is a copy of UCI ML housing dataset.

<https://archive.ics.uci.edu/ml/machine-learning-databases/housing/>

This dataset was taken from the StatLib library which is maintained at Carnegie Mellon University.

The Boston house-price data of Harrison, D. and Rubinfeld, D.L. 'Hedonic prices and the demand for clean air', J. Environ. Economics & Management, vol.5, 81-102, 1978. Used in Belsley, Kuh & Welsch, 'Regression diagnostics ...', Wiley, 1980. N.B. Various transformations are used in the table on pages 244-261 of the latter.

The Boston house-price data has been used in many machine learning papers that address regression problems.

.. topic:: References

- Belsley, Kuh & Welsch, 'Regression diagnostics: Identifying Influential Data and Sources of Collinearity', Wiley, 1980. 244-261.
- Quinlan, R. (1993). Combining Instance-Based and Model-Based Learning. In Proceedings on the Tenth International Conference of Machine Learning, 236-243, University of Massachusetts, Amherst. Morgan Kaufmann.

```
In [87]: data = pd.DataFrame(boston.data, columns = boston.feature_names)
```

```
In [88]: data.head()
```

Out[88]:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LSTAT
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	15.3	396.90	4.98
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	17.8	396.90	9.14
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	17.8	392.83	4.03
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	18.7	394.63	2.94
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	18.7	396.90	5.33

```
In [89]: data['Price'] = boston.target
```

```
In [90]: data.head()
```

Out[90]:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LSTAT	Price
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	15.3	396.90	4.98	24.0
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	17.8	396.90	9.14	21.6
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	17.8	392.83	4.03	34.7
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	18.7	394.63	2.94	33.4
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	18.7	396.90	5.33	36.2

```
In [91]: # See the column data types and non-missing values  
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 506 entries, 0 to 505  
Data columns (total 14 columns):  
CRIM      506 non-null float64  
ZN        506 non-null float64  
INDUS     506 non-null float64  
CHAS      506 non-null float64  
NOX       506 non-null float64  
RM        506 non-null float64  
AGE       506 non-null float64  
DIS       506 non-null float64  
RAD       506 non-null float64  
TAX       506 non-null float64  
PTRATIO   506 non-null float64  
B         506 non-null float64  
LSTAT     506 non-null float64  
Price     506 non-null float64  
dtypes: float64(14)  
memory usage: 55.4 KB
```

In [92]: *# Statistics for each column*
 data.describe()

Out[92]:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO
count	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000
mean	3.613524	11.363636	11.136779	0.069170	0.554695	6.284634	68.574901	3.795043	9.549407	408.237154	18.455534
std	8.601545	23.322453	6.860353	0.253994	0.115878	0.702617	28.148861	2.105710	8.707259	168.537116	2.164946
min	0.006320	0.000000	0.460000	0.000000	0.385000	3.561000	2.900000	1.129600	1.000000	187.000000	12.600000
25%	0.082045	0.000000	5.190000	0.000000	0.449000	5.885500	45.025000	2.100175	4.000000	279.000000	17.400000
50%	0.256510	0.000000	9.690000	0.000000	0.538000	6.208500	77.500000	3.207450	5.000000	330.000000	19.050000
75%	3.677083	12.500000	18.100000	0.000000	0.624000	6.623500	94.075000	5.188425	24.000000	666.000000	20.200000
max	88.976200	100.000000	27.740000	1.000000	0.871000	8.780000	100.000000	12.126500	24.000000	711.000000	22.000000

Exploratory Data Analysis

In [93]: *# Count the number of null values in the columns*
 data.isnull().sum(axis=0)


Out[93]:

CRIM	0
ZN	0
INDUS	0
CHAS	0
NOX	0
RM	0
AGE	0
DIS	0
RAD	0
TAX	0
PTRATIO	0
B	0
LSTAT	0
Price	0

dtype: int64

```
In [94]: # Visualize the Price distribution in our dataset.  
plt.figure(figsize=(10,5))  
plt.hist(data['Price'], color='red', bins=100)  
plt.xlabel('Price')  
plt.ylabel('Number of Homes')  
plt.title('Home Price Distribution')  
plt.show
```


```
Out[94]: <function matplotlib.pyplot.show(*args, **kw)>
```



In [95]: *# Visualize the Age distribution of houses in our dataset.*


```
plt.figure(figsize=(10,5))
plt.hist(data['AGE'], color='blue', bins=100)
plt.xlabel('Age')
plt.ylabel('Number of Homes')
plt.title('Distribution of Homes by Age')
plt.show
```

Out[95]: <function matplotlib.pyplot.show(*args, **kw)>





```
In [96]: # Visualise the Home Price distribution by Age
plt.figure(figsize=(10,5))
plt.scatter(x=data.AGE, y=data.Price, color='Red', alpha=0.5)
plt.xlabel('Age')
plt.ylabel('Price')
plt.title('Home Prices by Age')
plt.show
```

```
Out[96]: <function matplotlib.pyplot.show(*args, **kw)>
```




```
In [97]: # Visualise the Home Price distribution by number of Rooms
plt.figure(figsize=(10,5))
plt.scatter(x=data.RM, y=data.Price, color='green', alpha=0.5)
plt.xlabel('Number of Rooms')
plt.ylabel('Price')
plt.title('Home Prices by Number of Rooms')
plt.show
```

```
Out[97]: <function matplotlib.pyplot.show(*args, **kw)>
```



```
In [98]: # Visualize the Room distribution in our dataset.  
plt.figure(figsize=(10,5))  
plt.hist(data['RM'], color='red', bins=100)  
plt.xlabel('Room')  
plt.ylabel('Number of Homes')  
plt.title('Distribution of Rooms')  
plt.show
```

```
Out[98]: <function matplotlib.pyplot.show(*args, **kw)>
```



```
In [99]: # Seems like there is a one home with a large number of rooms priced pretty low. Lets see if this is an outlier  
data['RM'].describe()
```

```
Out[99]: count      506.000000  
mean         6.284634  
std          0.702617  
min          3.561000  
25%          5.885500  
50%          6.208500  
75%          6.623500  
max          8.780000  
Name: RM, dtype: float64
```

```
In [100]: data['RM'].sort_values().tail(15)
```

```
Out[100]: 283      7.923  
166      7.929  
204      8.034  
226      8.040  
97       8.069  
233      8.247  
253      8.259  
224      8.266  
267      8.297  
232      8.337  
163      8.375  
262      8.398  
257      8.704  
225      8.725  
364      8.780  
Name: RM, dtype: float64
```

```
In [101]: # Look at the prices based on the number of rooms for the records above
data.loc[[97,163,224,225,232,257,267,364]]
```

Out[101]:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LSTAT	Price
97	0.12083	0.0	2.89	0.0	0.445	8.069	76.0	3.4952	2.0	276.0	18.0	396.90	4.21	38.7
163	1.51902	0.0	19.58	1.0	0.605	8.375	93.9	2.1620	5.0	403.0	14.7	388.45	3.32	50.0
224	0.31533	0.0	6.20	0.0	0.504	8.266	78.3	2.8944	8.0	307.0	17.4	385.05	4.14	44.8
225	0.52693	0.0	6.20	0.0	0.504	8.725	83.0	2.8944	8.0	307.0	17.4	382.00	4.63	50.0
232	0.57529	0.0	6.20	0.0	0.507	8.337	73.3	3.8384	8.0	307.0	17.4	385.91	2.47	41.7
257	0.61154	20.0	3.97	0.0	0.647	8.704	86.9	1.8010	5.0	264.0	13.0	389.70	5.12	50.0
267	0.57834	20.0	3.97	0.0	0.575	8.297	67.0	2.4216	5.0	264.0	13.0	384.54	7.44	50.0
364	3.47428	0.0	18.10	1.0	0.718	8.780	82.9	1.9047	24.0	666.0	20.2	354.55	5.29	21.9

```
In [102]: # Based on the comparison with other columns in the dataset, Does not seem like it is an outlier.
# Therefore I will keep the data for that home.
```

Since my dataset has no null values and all are numeric values I will start working on the ML Model

Building a Random Forest Model

Split the dataset into training set and test set

```
In [119]: # Split the feature set and target set

features = data.drop(columns='Price')
targets = data['Price']
```

In [120]: *# Split the data into 75% training set and 25% test set*

```
X_train, X_test, y_train, y_test = train_test_split(features, targets, test_size = 0.25, random_state = 0)

print(X_train.shape)
print(X_test.shape)
print(y_train.shape)
print(y_test.shape)
```

```
(379, 13)
(127, 13)
(379,)
(127,)
```

In [121]: *# Since the dataset is in different units (Age = Years, Distance = Miles, Tax = Dollars),
apply standard scaler to the X_train and X_test.*

```
sc = StandardScaler()

# Transform and fit both the training and testing data

X_train = sc.fit_transform(X_train)
X_test = sc.fit_transform(X_test)
```

In [122]: *# Create the model*

```
rfg = RandomForestRegressor(n_estimators= 30, random_state=0)
rfg.fit(X_train, y_train)
```

Out[122]: RandomForestRegressor(bootstrap=True, criterion='mse', max_depth=None,
max_features='auto', max_leaf_nodes=None,
min_impurity_decrease=0.0, min_impurity_split=None,
min_samples_leaf=1, min_samples_split=2,
min_weight_fraction_leaf=0.0, n_estimators=30, n_jobs=None,
oob_score=False, random_state=0, verbose=0, warm_start=False)


In [109]: *# Make predictions*

```
rfg_pred = rfg.predict(X_test)
```

In [142]: *# Visualize the Actual Price compared to the predicted price.*

```
plt.figure(figsize=(8,8))
plt.scatter(y_test,rfg_pred, color = 'green')
plt.xlabel('Actual Prices')
plt.ylabel('Predicted Prices')
plt.title('Actual Prices vs Prices Predicted')
plt.show
```

Out[142]: <function matplotlib.pyplot.show(*args, **kw)>



In [137]: *# Evaluate the model performance metric*

```
score = r2_score(y_test, rfg_pred)
print('R Squared Score of the Test data is: R^2 = %0.5f' % score)
print('Random Forest Regression ML Model Performance on the Test data is: MAE = %0.5f' % mean_absolute_error(y_test, rfg_pred))
```

R Squared Score of the Test data is: R^2 = 0.75635

Random Forest Regression ML Model Performance on the Test data is: MAE = 2.91562

In [138]: *# Using the given Boston Housing data, our machine Learning model can predict the house prices within 3 points
at a 76% accuracy.*

In []: